

# Exercise 1 Delivery: Constructive Analysis

Fellipe Pessanha

September 19, 2025

## 1 Exercise Description

The task in hand was to implement solution constructors for the specified software dependency problem, in specific a *random constructor*, *randomized greedy constructor*, and a *greedy constructor*, then perform a statistical analysis of their performance.

## 2 Implementation Details

The implementation was carried out using **Julia**. The implementation steps involved

1. Parsing the input data to extract problem context
2. Implement a prototype evaluator in order to assess solution quality
3. Write tests that make sure each step is working as expected

The best way to understand the code structure is to check the repository's test suite at `./test/runtest.jl`

### 2.1 Constructive implementation

The randomized greedy constructive was implemented using a combination of a max heap  $H$  – keeping track of all the available packages sorted by their *benefit* – and a vector  $V$  – used to store all the  $n$  best evaluated packages by the same criterion, guaranteeing optimal  $O(N \cdot \log(\#H))$ , where  $N$  is the number of packages selected packages, and  $\#H$  is the mean number of packages available in the unused heap throughout the execution.

The size of the vector is determined by the parameter  $\alpha \in [0, 1]$ , and the random constructive was simply implemented by setting  $\alpha = 1$ , and the greedy constructive by setting  $\alpha = 0$ .

## 3 Analysis Methodology

The analysis was performed for 30 equidistant values of  $\alpha \in [0, 1]$ .

In each step, 30 independent runs were performed for each constructor, and captured the average solution quality, standard deviation, and execution time, evaluated with Julia's BenchmarkTools package, which accounts for garbage collection, precompilation time, and other Julia specific nuances.

## 4 Results

The results of the analysis, normalized in % values, can be seen in Figure 1.




Figure 1: mean solution quality, standard deviation, and time elapsed by  $\alpha$

We can see in our plots that the execution time, aside from some noise and outliers, is mostly consistent during the entire range of  $\alpha$  values, with a slight trend upwards as  $\alpha$  increases. The standard deviation also seems mostly regular for  $\alpha \in [0.1, 1]$ , and the mean value of the solutions consistently decreases as  $\alpha$  increases, with the purely random constructor ( $\alpha = 1$ ) performing at about 50% of the greedy constructor ( $\alpha = 0$ ).

## 5 Conclusion

With the results in hand, we can conclude that, outside of a small range of small  $\alpha$ , our constructor behaves consistently, having a standard deviation of a mostly random sample of values.

Knowing that  $\alpha$  has very little impact on the execution time, we can also conclude that the choice of  $\alpha$  can be manipulated to achieve the desired balance of solution quality and variety, and that values with  $\alpha \in [0.3, 0.6]$  seem to be a good sweet spot.